

## Poll() function

"It is a system call that select devices / files from the list of available files / devices that are ready for intended operation."

→ Select() is used for checking device for read, write and error, Poll() is extended version of select which also check a device is ready for block read, normal read, character read etc.

It has more flags than select.

```
int poll(struct pollfd poll_fd1[], int n,  
         int timeout)
```

→ Poll fd1[] has list of specific devices to be checked, select generally take a list of devices and check.

① Struct poll fd

```
{  
    int fd;  
    int events; → purpose  
    int revents; → response  
}
```

② poll fd1[];

→ list of fds to be checked, each device / file has one poll fd.

③ int n;

Total number of devices / files

④ int timeouts;  
→ timeout in seconds.

Q Write a program for two files, bluetooth and printer to check for read and write, will check ~~1~~ 2 read operations and 1 write operation.

Bluetooth → Read and Write  
Printer → read

Total devices = 2, operations ≈ 3

#include <poll.h>

int main()

{

char buf[1000];

int printer\_fd = open("/dev/printer",  
O\_WRONLY);

int bt\_fd = open("/dev/bluetooth",  
O\_RDWR);

struct pollfd pollfd1[2];

pollfd1[0].fd = printer\_fd;

pollfd1[0].events = (pollfd1[0].events | POLLRDNORM | POLLWRNORM);

pollfd1[1].fd = bt\_fd;

pollfd1[1].events = (pollfd1[1].events | POLLRDNORM | POLLWRNORM);

int nfds = poll(pollfd1, 2, 30);

int (nfds = 0)

{

```
    prntf ("No device ready ");  
}  
if (nfdi = -1)  
{  
    perror (errno);  
}  
if (nfdi > 0)  
{  
    if (poll_fd1[0].revent & POLLWRNORM)  
    {  
        write (pwriter_fd, " HELLO", 5);  
    }  
    if (poll_fd1[1].revent & POLLRDNORM)  
    {  
        read (bt_fd, buf, 1000);  
    }  
    if (poll_fd1[1].revent & POLLWRNORM)  
    {  
        write (bt_fd, "Hello", 5);  
    }  
}  
close (pwriter_fd);  
close (bt_fd);  
return 0;  
}.
```

## File Representation

### Process 1

```
int fd1 = open("file.txt", O_RDONLY);
int fd2 = open("file.txt", O_WRONLY);
```

FDT - P1

STDIN	0
STDOUT	1
STDERR	2
fd 1	3
fd 2	4

- Whenever a file is opened an entry is created in the system file table and file descriptor table.

System file table

file
read, offset = 0 counter = 1
write, file_offset = 0 counter = 1
write, file_offset = 0 counter = 1

inode table.

file1.txt, disc_id;
disc_serial, disc_size;
created_user, created_group
permissions, creation_time
last_access_time, last_modification, inumber = 1
file1.txt discid -----

counter → tells that how many pointers are pointing towards the file.  
 file\_offset → It is a pointer which points that how many characters have been read / write.

## Process 2

```
int fd3 = open("file.txt", O_WRONLY);
```

FDT - P2

STDIN	0
STDOUT	1
STDERR	2
fd3	= 3

- An entry is created in the system file table.
- If there is already an entry in the inode table then no. new entry will be created. and file from the system file table points toward the same location.
- Pointer updated
- At a time there will be only one entry of a file in inode table.
- When parent has opened a file and fork is called then the child has access to that file as file descriptor table has been copied.
- Process 2 has created a child.

FDT - Ch - P2

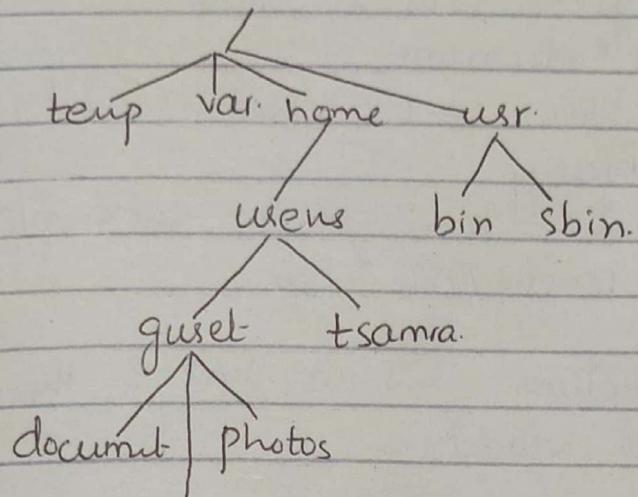
STDIN	0
STDOUT	1
STDERR	2
3	(fd, 1)

- Counter will be updated from 1 to 2 in fd1 and fd2 in Fd table.
- For each call of close system call → one entry is deleted from the "system file table" and one or more from "file descriptor table".
- For each close system call an entry is deleted from inode table if entry is 1, otherwise counter is decremented.

## Chapter #05

### File and Directories.

"/" is the home directory  
in this we have other directories



Downloads.

In unix system directory is actually a file.

In simple file we keep content but in directory file it consist of the other files stored in it.

"." is current working directory  
".." is the previous directory.

→ current working directory is accessed by (pwd).

int getcwd (char \*buf, int size)

↓  
Where to put  
That current working  
directory

int  
fcntl

→ change current working directory (cd)  
`chdir (char * dirname);`

→ directory oper.

For this system call is opendir

→ `DIR * opendir (char * dirname)`

For DIR return type we should include  
`<dirent.h>` library.

Now if i want to open the same file.

→ Struct dirent \* readdir(DIR \* dir);

dirent in a structure and it has 3 things

→ size of file or directory

struct dirent

```
{ d_size // size; ✓  
  d_name // name;  
  d_ctime // creation time✓  
}
```

int close dir (DIR \* dir )

→ Program to implement ls - command.  
report number of sub-directories, count

Steps

- ① create/declare DIR and dirent pointers.
- ② open the given directory.
- ③ read the given from directory and list until we reach NULL.
- ④ For each entry list and find its type
- ⑤ if type is directory update count.
- ⑥ Print +.

*short*  
close directory.

### Program

```
#include <dirent.h>
int main( int argc, char *argv[] )
{
    DIR *dir; // point ①
    struct dirent *dirp; // point ①
    int dcount=0;

    if( dir = opendir( argv[1] == NULL ) // point ②
        { perror("Error in opening of dir");
          exit(0); }
    else
        while( (dirp = readdir( dir )) != NULL )
        {
            printf("%s", dirp->d_name);
            if( isDirectory( dirp->d_name ) )
                dcount++; // it return true
                update counter.
        }

    printf("Number of subdirectories: %d\n",
          dcount);

    closedir( dir );
    return 0;
}
```

which command -?  
echo command -?

## File State Information

If we know the name of file or its hard link or we had opened it and we had its file descriptor and want to extract the information of that file in inode file, so we will use certain system calls() for it.

e.g. which device, its size, its user etc type of information.

The systemcalls used for this are:

stat(char\* path, struct stat \*buf)

lstat(char\* link\_path, struct stat \*buf)

fstat(int fd, struct stat \*buf) // we pass the file descriptor

The return type of all the three is integer.  
if failed returns -1.

Struct stat {

dev\_t st\_dev; // device id

ino\_t st\_ino; // what is inode number

mode\_t st\_mode; // permissions of file.

uid\_t st\_uid; // created by user.

gid\_t st\_gid; // user group

off\_t st\_off; // size of file.

nlink\_t st\_nlink; // hard links

time\_t st\_atime; // last access time

```
time_t st_ctime; // correct time  
time_t st_mtime; // modification time  
}
```

→ Program to find that which of the two files is greater/larger.

```
#include <sys/stat.h>  
int main()  
{  
    struct stat *file1_buf = (stat*) malloc(sizeof(stat));  
    struct stat *file2_buf = (Stat*) malloc(sizeof(Stat));  
  
    stat("file1.txt", file1_buf);  
    stat("file2.txt", file2_buf);  
  
    if (file1_buf->st_size > file2_buf->st_size)  
        printf("file1 is larger");  
    else if (file2_buf->st_size > file1_buf->st_size)  
        printf("file2 is larger");  
    else  
        printf("Both are equal size");  
    return 0;  
}
```

⇒ In Linux we had both soft-links and hard-links.  
 relation with inode table:

hard-link points to the  
 inode number of that file.  
 soft-link points to the name of that file.

If I change the name of file the hard-link to it will not break but the soft-link to it will break.

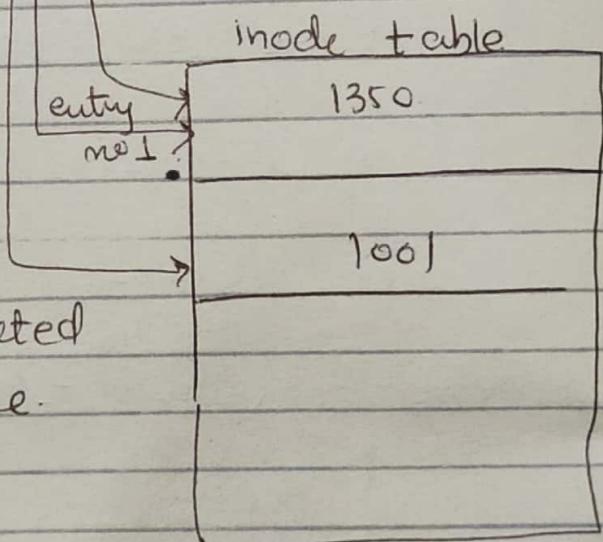
If we delete a file, it not actually deleted from the disk only its data is deleted from the inode-table.

### hard link

```
int link (char* path1, char* path2);
```

Original file name & path.

dir 1		
file	file	inode no
file1.txt	1350	
mycode.c	1001	
link_file1.txt	1350	



```
int symlink (char* path1, char* path2);
```

name and path of hard link to be created.

pointer to - - h

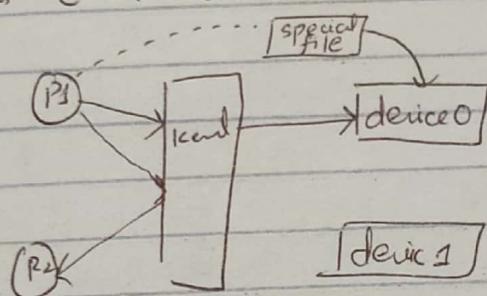
int unlink (char \* path);

x ————— x ————— x ————— x ————— x

Chapter # 06

File Pipes (Interprocess communication).

let we have two process P1 and P2 and two devices d0 and d1. Now these process can access through kernel



P1 access to device0 and open it, when we open a device actually a special file is opened.

Now if the two process wants to communicate through each other e.g. to send data it will be done through interprocess communication using kernel.

One method is two used socket method in this we must have a network card on our computer.

2nd method is to use shared libraries.

Another method used by UNIX is "pipes".

In this type if two process wants to communicate kernel creates a "special file" one process reads to it and other writes to it in this way data is shared.

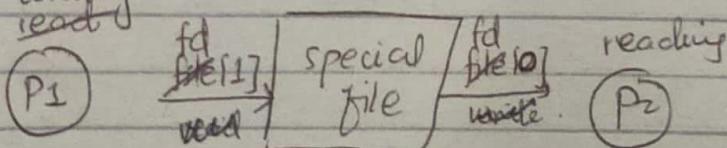
But this shared file is always created by kernel so the process should send a request through a system call.

`int pipe(int fd[])` → // sending request to kernel to start I want to get a pipe IPC.

As a result of this call, kernel creates a special file.

Now how reading writing will be done in special file so the

`int fd[2]` array of size 2 will be used to read all write a data by making handles



### Disadvantage:

The disadvantage is that we cannot know its name so we cannot use it for another process except for the child process.

Program to create a process which will create a pipe will write the data and then will create a child process and the child will display the parent data.

### Points

- create int array size 2.
- create pipe.
- write "Hello" via fd[1].  
to the pipe.
- create child process.
- Child process read data from pipe via via fd[0].
- child process display the read data.

library used is `unistd.h`.

```
#include <unistd.h>
```

```
int main()
```

```
{
```

```
    int fd[2];
```

```
    int pipe((fd, Fd));
```

```
    write(fd[1], "Hello from your parent process",
```

```
        int pid = fork();
```

```
        char * buf[29];
```

```
        29);
```

```
        // to read the data in
```

```
        buffer.
```

```
    read(fd[0], buf, 29);
```

```
    printf("%s", buf);
```

```
    return 0;
```

Another System Call used for I/O

FIFOs (named pipes)

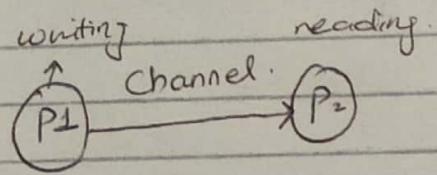
`mkfifo(char * pipename, int mode)`

Loggers are the separate applications in the system and we can communicate through it using `mkfifo`.

For using this channels (special files) will be created one for reading and one for writing.

In previous the special files were opened by the kernel.  
But in this we had to open them.

Program to write so that two processes communicate using fifos.



For P1.

```
#include <unistd.h>
```

```
int main()
```

```
{  
    int mkfifo ("chat1", S_IRUSR | S_IWUSR);  
    int fd = open ("chat1", O_WRONLY);  
    write (fd, "Hello", 5);  
    write (fd, "How are u doing?", 12);  
    return 0;  
}
```

For P<sub>2</sub> reads

```
#include <unistd.h>
int main()
{
    int fd = open("chat1", O_RDONLY);
    char buf[50];
    read(fd, buf, 50);
    printf("%s", buf);
    return 0;
}
```

I RATHER

### \* Pipeline:

lets you have one terminal and you run commands.

lets from a file i want to know that how many are there in a file.

so we will use

\$ grep error output

but this command is not enough we want to know the number of lines to be created.

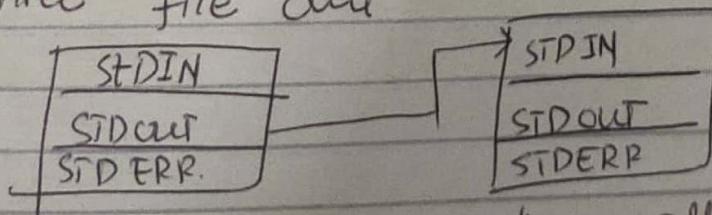
(P<sub>1</sub>)

(P<sub>2</sub>)

\$ grep error output | wc -l

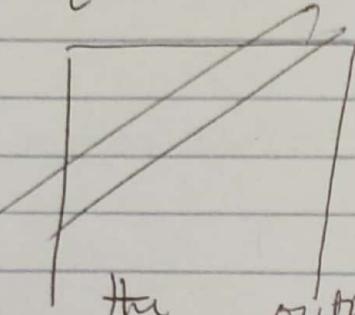
this is pipeline.

for every process there are at least three file des



but if we pipeline all the data

and commands of  $P_1$  will be transferred to  $P_2$ .



so the output of one process will become the input of another process.

## Chapter # 08

### Signals

Since them are interrupt that are always directed towards the CPU.

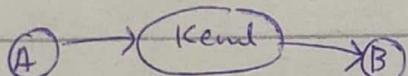
But signals are these short notification or messages send to the process as a result of an event.

they are generally generated from one process to another but kernel is involved in it.

e.g sleep(10)

send a signal to kernel to suspend the process of a process for 10 seconds.

signal from A to B through kernel.



SIGABRT → About the process → implementation

SIGBUS → Access invalid memory → implementation

dependent dependent.

SIGFPE → Division by zero → "

SIGKILL → Kill the process → not implementation

SIGALRM → we pass time in it, when the time is complemented it kill the process.

SIGCHLD → on child termination → implementation dependent.

SIGINT → Abort → implementation dependent.

SIGUSR1 → unused

SIGUSR2 → unused.

SIGSTOR → stop execution → not implementation dependent.

How signals are generated with  
shortcuts like  $\text{ctrl}+\text{x}$ ,  $\text{ctrl}+\text{c}$   
it is told by a command.

$\$ \text{stty}-a \rightarrow$  list of signals that are  
generated with shortcuts.

$\$ \text{kill}-l$  → gives you the list of all  
signals generated

⇒ How to send a signal from terminal  
to a process?

$\$ \text{kill}-s$  signal name pid → send signal  
to process.  
e.g. SIGABRT 0034

A process can receive a signal only  
if it is in execution state on  
the CPU.

If the process is not in execution  
state the signal should wait for it.

A signal is delivered only when the  
process is in execution.

But the signal delivered will be in  
ordered like the signal delivered  
first will be received by the process  
first.

⇒ If a process wants to send  
signal to another process.

int kill(pid\_t pid, int signo); → definition  
kill(3256, SIGINT); call.

`kill(getppid(), SIGKILL);`

⇒ If a process wants to sent signal to itself:

```
int raise(int sig_no);
raise(SIGKILL);
```

⇒ If we want to delay a signal this is known as 'blocking signal'.  
if we want to block a signal we had to made a list of these signal which we want to block.

`sigset_t set;`

⇒ `int sigaddset(sigset_t *set, int signo);`  
⇒ if we want to delete a signal from set.

↓  
⇒ `int sigdelset(sigset_t *set, int sig_no);`

⇒ ~~int emptyset(sigset\_t \*set);~~  
⇒ `int sigemptyset(sigset_t *set);`

⇒ `int sigfullset(sigset_t *set);`

⇒ `int sigismember(sigset_t *set, int sig_no);`

⇒ header file for all of them is `<signal.h>`.

Now if we want to do the above tasks with the system call.

`sigprocmask(int how, sigset_t *newset, sigset_t *oldset)`

syscall used to block or unblock signals or setup signals  
3 options  
SIG\_BLOCK  
SIG\_UNBLOCK  
SIG\_BLOCK

If you block a signal in the above argument tell us the signals that were blocked previously

## Lecture

### Handling of Signals

`sigaction(int sig_no, struct sigaction* act,  
              struct sigaction* oact)`

`sigaction* act` → action to be applied on  
the reception of the signal.

There is a signal and you want to  
say the signal that i want to change  
the default action of that signal.

`sigaction* oact` → old action of the signal.

We use that oact that if we restore the  
signal action back, we should have  
it.

`struct sigaction {`

```
    void (sa_handler * )int; // handler 1  
    sigset * sa_mask;  
    int sa_Flags;
```

void (sa\_sigaction\_handler)(int, sigset \*, void)  
  // handler 2.  
  Real time.  
}

Function of all these.

- ① void (sa\_handler \*) int // normal POSIX define Function  
the handler you passed in this will be  
given as an action to the signal.

e.g

DEL            // give default action.  
IGN            // ignore signal.  
handler function // user define function  
  which we pass in  
  it.

- ② sigset \* ~~same as~~:

40 signals to be blocked during process  
i.e while handling some signals we  
want some other coming signals to  
be blocked, so we will need it.

- ③ int sa\_flags:

// select handler function:  
Here we have two type of handles  
by passing 0 and 1 in flags we  
can select either of them.  
about X

SIGFPE → division by zero  
passing this signal need the signal  
to be aborted so we want  
to change its default action.

```
#include <signal.h>
int main()
{
    struct sigaction act, oact;
    sigset(SIGFPE, FPE_mask);
    // Since we want to block two
    // signals so we will add them
    // to set
    sigaddset(&FPE_mask, SIGINT);
    sigaddset(&FPE_mask, SIGUSR1);
    act.sa_handler = FPE_signal_handler;
    act.sa_mask = FPE_mask;
    act.sa_flags = 0;
    if (sigaction(SIGFPE, &act, &oact) != 0)
        return 0;
}
```

```
// function:
void FPE_signal_handler(int sig_no)
{
    printf("Division by zero was attempted. I am
           avoiding the about process using
           custom handler.");
}
```

## Waiting For a Signal

e.g We are sleeps to sleeps a signal for a certain time.

int sleep(int seconds);

→ int pause(void); → suspend the process, until a signal is received.

The signal received (tak hm ko need se jaga) we dont know its number.

so what we do is:

→ int sigsuspend(sigset \*set1);

We will a set of 3 signals juk hmen jagen gyj.

→ int sigwait(sigset \*set1, int \*sig\_no)

We will use a set of signals but we will also have its number.

## Chapter # 09

### Times and Timers.

Time is a running entity, it increments.  
e.g. we want to note the time b/w  
two events or we want to know that  
how much time a process takes:  
`<time.h>`

lets me a syscall:

- ① `time_t time(time_t * t1)` → return the time  
~~epoch~~ elapsed since epoch.  
→ epoch is actually the time after which  
we want to know that how much  
seconds passed.  
e.g. Jan 01, 1970, 12AM. UTC.

- ② ~~double~~ `double difftime(time_t t1, time_t t2);`  
it gives the time difference b/w two  
functions as process  
int is of 4 byte  
long int is of 8 byte.  
double is of 10 byte.

The issue here is that the above function  
give us time in seconds, with no year,  
no day,

- ③ `char* ctime(time_t t1)`  
we passed the time returned by  
syscall 1 to it and it return the  
time into string format.

A better syscall than this is

struct tm\* localtime (time\_t t1);  
convert time into tm structure

```
struct tm {  
    int tm_sec;  
    int tm_min;  
    int tm_hour;  
    int tm_year;  
    int tm_mday;  
    int tm_mon; // month  
    int tm_wday; // weekday  
    int tm_yday;  
    int tm_isdst; // is day light saving  
};  
time on or off.
```

wday means e.g. since monday how many days passed.

④ struct \* tm\* gmtime (time\_t t1);  
// convert time into structure (UTC).

⑤ char\* asctime (struct time\_\* tm1);

// convert time from from tm into string.

CPU can execute 1 billion instruction at one second, so if the process takes time less than 1 second so the above ~~instruction~~ syscall return 0. To solve this problem, we have another syscall

⑥ `int gettimeofday(struct timeval* tv, void* tap);`

```
struct timeval {  
    time_t tv_sec;  
    time_t tv_usec;  
};
```

returns time since epoch in seconds and microseconds.

⇒ To know about the execution time of two processes:

```
int main()
```

```
{  
    time_t t1, struct timeval * t1, t2;  
    int gettimeofday (tv1, NULL);  
    // function whose time we want to know:  
    int funt = factorial (15);  
    gettimeofday (tv2, NULL);
```

long int execution\_time\_in\_sec = ~~tv2 - tv1~~

$(tv_2 \rightarrow tv\_sec - tv_1 \rightarrow tv\_sec) *$

$1000000 + (tv_2 \rightarrow usec - tv_1 \rightarrow usec)$

$tr\_$

upto how we are taking time since epoch  
it was giving the time of execution.  
When we run process on the CPU.

For sometime it runs on the CPU and  
for sometime it is removed from the CPU.  
So the time of execution of a function  
differs from computer to computer.

Now how much time a process  
have taken on CPU is found with  
syscall

```
clock_t time ( struct tms * t1 )
```

```
struct tms {
```

```
clock_t
```

tms\_utime; <sup>①</sup> // user process time spent of user code

```
clock_t
```

tms\_stime; <sup>②</sup> // system time on behalf of user process

```
clock_t
```

tms\_cutime; // user process time terminated child

```
clock_t
```

tms\_cstime; // system time co behalf of  
process and terminated children )

```
}
```

So when we use this the ① and ② will tell that CPU will tell that this much time is take by your code and this much time i had worked from you.

(i) In system we also have

⇒ REAL TIME CLOCK:

This starts when our system starts and this gives the time since epoch in seconds and nano seconds. It has three system calls.

- ① int clock\_gettime(clockid\_t clockid, struct timespec\* res);
- ② int clock\_gettime(clockid\_t clockid, struct timespec\* res);
- ③ int clock\_settime(clockid\_t clockid, struct timespec\* res);

Struct timespec

time\_t tv\_sec

long tv\_nsec;

}

⇒ gettime() gives us the time from which the system has been started. i.e

⇒ getres() through this we reset the clock. Usually the clock starts itself and then increments it but the user can reset it according to its needs.

⇒ clock\_settime is used e.g. if tell the system that I don't know from which time clock is started, but I give a time in clock\_settime() function, which starts that starts the function clock from that given time but for this we also need a root user.

Now we will go to two other calls

① int sleep(time\_t t);  
the time given in it will make a signal sleep upto that time.

② int nanosleep (struct timespec \*tm, struct timespec \*rtm).

It is a modified version of sleep.  
it can take the time in second as well  
as in nano-second. and rtm is the  
resolution time.

Timers:-

one important concept is of stopwatch  
i-e we send it certain time.

→ int gettimer(int which, struct itimerval \* value);

→ int settimer(int which, struct itimerval \* value,  
                  struct itimerval \* oval)

Struct itimerval

{

    struct timeval value;  
    struct timval interval;

}

There are three type of timers.

① ITIMER\_REAL; //decrements according to real time clock  
ITIMER\_VIRTUAL; //decrements according to process time on CPU and generate SIGALRM when timer expires.  
ITIMER\_PROF//decrements according to process time and kernel time and generates SIGPROF when timer expires.

6  
the int which parameter in the above function take any of these three time that which timer do you want.

struct timeval

```
{ time_t tv_sec;  
  time_t tv_usec;
```

}  
then if you give it a counter time you will pass struct timeval value;

→ then it asks that if this timer expires what to do?

There will be two possibilities either to send the alarm and then the alarm will take some further time

after that again ALARM will ring.

For this we use struct timered (interval for reset).

Write a C system Program that creates a child process everytime after 30 sec.

```
#include <time.h>
int main()
{
    struct sigaction *act;
    sigsetmask(SIG_BLOCK(SIGUSR1));
    struct itimerval;
    struct timeval;
    act.sa_handler = sigthandler();
    sigaddset(SIG_BLOCK(SIGUSR1));
    act.sa_mask = mask;
    act.sa_flag = 0;
    sigaction(SIGALRM, &act, NULL);
    it.value.tv_sec = 30;
    it.value.tv_usec = 0;
    it.interval.tv_sec = 30;
    it.interval.tv_usec = 0;
    setitimer(ITIMER_VIRTUAL, &it, NULL);
    return 0;
}
void sigthandler(int sig_no)
{
    fork();
}
```